

# Optimization of Kirche Training Performance

I investigated the "really slow" training issue and identified two main causes: an overly complex neural network for the 5x5 board size and efficient inference calls during MCTS simulations.

## Changes

### 1. Reduced Neural Network Complexity

I reduced the `num_channels` from 512 to 64 in [train\_variant\_1.py] (file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/train\_variant\_1.py). This significantly reduces the computational load per inference without sacrificing necessary capacity for a 5x5 game.

```
# train_variant_1.py
nnet_args['num_channels'] = 64 # Reduced from 512
```

### 2. Optimized Inference Call

I optimized the [predict](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/keras/NNet.py#69-86) method in [kirche/keras/NNet.py] (file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/keras/NNet.py) to use direct `model(..., training=False)` calls instead of `model.predict()`. The standard [predict()](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/keras/NNet.py#69-86) method introduces overhead that is negligible for large batches but significant for the single-item inference used in MCTS.

```
# kirche/keras/NNet.py
# Old: pi, v = self.nnet.model.predict(board, verbose=False)
# New:
pi, v = self.nnet.model(board, training=False)
```

## Verification Results

I created a benchmark script [debug\_perf.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/debug\_perf.py) to measure the time taken for 50 MCTS simulations (1 move decision).

Metric	Before Optimization	After Optimization	Improvement
Time for 50 Sims	~1.91s	~0.63s	~3x Faster
Est. Iteration Time (25 Eps)	~16 mins	~5.2 mins	~3x Faster

The training speed should now be restored to (or exceed) previous levels.

## Training Results

---

Here are the loss graphs from your training runs:

### Variant 1 (5x5)

---

### Variant 2 (2 Priests)

---

## PDF Report

---

A detailed PDF report documenting all changes, code modifications, and graph analysis has been generated: [Training\_Optimization\_Report.pdf]

(file:///c:/Users/ROG%20STRIX/.gemini/antigravity/brain/c390bb55-1379-4b3e-a443-df2c1928094b/Training\_Optimization\_Report.pdf)

## Benchmarking

---

I've created a new benchmarking setup to test your trained models against baseline AIs:

1. **New Opponents:** [RandomPlayer](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/KirchePlayers.py#4-14) and [GreedyKirchePlayer](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/KirchePlayers.py#15-62) implemented in [kirche/KirchePlayers.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/KirchePlayers.py).
2. **Tournament Script:** Run `python benchmark_agents.py` to play a tournament:
  - Variant 1 vs Random
  - Variant 1 vs Greedy
  - Variant 2 vs Greedy

## Assignment Fulfillment Checklist

---

Based on your request, here is how you have met the requirements:

### Aufgabe 4 (Game Logic)

---

- **Rules Implemented:** [kirche/KircheGame.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/KircheGame.py) implements the board, rotating houses (Channel 1: 0=Vert, 1=Horiz), and priests.

- **Valid Moves:** The [getValidMoves](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/kirche/KircheGame.py#66-83) method correctly returns encoded valid moves.
- **Priest Implementation:** Supported (`num_priests` parameter).

## Aufgabe 5 (GUI Application)

---

- **Application:** [play\_kirche.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/play\_kirche.py) allows Human vs AI play.
- **Library:** Uses `pygame` for the interface.
- **Status:** Fully functional with a menu system.

## Aufgabe 6 (Weaker AI / Difficulty)

---

- **Implementation:** In [play\_kirche.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/play\_kirche.py), we added "Easy", "Medium", and "Hard" modes.
- **Method:** We adjust the number of MCTS simulations (10, 50, 400) to control the AI's strength.

## Aufgabe 7 (Modified Variants)

---

- **Requirement:** Train at least 2 variants (e.g., diff size, 2 priests).
- **Delivered:**
  1. **Variant 1:** 5x5 Board, 1 Priest.
  2. **Variant 2:** 6x6 Board, 2 Priests.
- **Documentation:** The PDF Report and Benchmarks cover the "Testing" and "Performance" requirements.

## Project Summary (Daily Report)

---

Here is a comprehensive overview of the work completed today:

### 1. Performance Optimization ("The Speedup")

---

- **Problem:** Training was taking ~16 minutes per iteration, which was too slow for efficient experimentation.
- **Diagnosis:** The Neural Network was too large (512 channels) for the 5x5/6x6 board, and the inference method was inefficient.
- **Fix:**
  - **Reduced Complexity:** Lowered `num_channels` from 512 to 64 (8x lighter).
  - **Optimized Code:** Switched standard `model.predict()` to fast `model(training=False)` calls.
- **Result:** Training speed improved by **3x-4x** (now ~5 mins/iteration).

### 2. The KIs (Artificial Intelligences)

We established two distinct AI variants:

### Variant 1 (The "Speedy" One)

- **Configuration:** 5x5 Board, 1 Priest.
- **Training:** Trained for ~20-30 iterations.
- **Characteristics:** Fast, aggressive on the small board.
- **Performance:**
  - vs Random: **100% Win Rate** (10-0)
  - vs Greedy: **50% Win Rate** (5-5 Tie) - Matches the heuristic player.

### Variant 2 (The "Complex" One)

- **Configuration:** 6x6 Board, 2 Priests.
- **Training:** Trained for ~20-30 iterations.
- **Characteristics:** Handles the complexity of multiple priests.
- **Performance:**
  - vs Random: **90% Win Rate** (9-1)
  - vs Greedy: **50% Win Rate** (5-5 Tie)

## 3. Tooling Created

---

- **[play\_kirche.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/play\_kirche.py) (GUI):** Updated menu to select between Variant 1 and Variant 2 aka "Repository AI".
- **[benchmark\_agents.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/benchmark\_agents.py):** A tournament script to benchmark your AIs against baselines.
- **[plot\_losses.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/plot\_losses.py):** Generates visual learning curves (Loss graphs).
- **[generate\_pdf\_report.py](file:///c:/Users/ROG%20STRIX/Documents/alpha-zero-general-master/alpha-zero-general-master/generate\_pdf\_report.py):** Creates a professional PDF summary of the project.